

# YASHWI PASSARY

+1 (812) 778-5805 | [ypassary.masters@gmail.com](mailto:ypassary.masters@gmail.com) | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

## SUMMARY

Customer-facing Data and AI Engineer with 5+ years of experience delivering production systems across healthcare, pharma and financial services. Built Python, SQL and Spark solutions on AWS, Azure, Databricks, Snowflake and Palantir Foundry, including 2B+ record migrations and operational workflows supporting 600K+ patients. Combines hands-on engineering with stakeholder-led delivery, UAT, production rollout and measurable business impact.

## EDUCATION

**M.S., Data Science** | *Indiana University – Bloomington*

**Aug 2024 - May 2026**

**B.Tech, Computer Science** | *Netaji Subhash Engineering College*

**Aug 2016 - Jul 2020**

## PROFESSIONAL EXPERIENCE

**Henry Ford Health**

**Jun 2026 - Present**

*Senior Data Engineer*

*Detroit, Michigan*

- Automated daily payer and insurance claims data pipelines using Fabric and SSMS, converting complex vendor data into contract classifications, risk categories, and Power BI insights for over 600K members, reducing manual processing time and supporting ongoing data growth.
- Built, validated, and deployed production reporting solutions using Epic Clarity and Caboodle data for Medicare, Medicaid and MMP data, enabling reliable member and claims-level reporting for payer and operational stakeholders.

**Project 990**

**Aug 2025 - May 2026**

*Data Scientist*

*Remote, USA*

- Prepared an entity matching pipeline resolving 100K+ duplicate nonprofit records using TF-IDF, cosine similarity, phonetic matching, and geographical features, achieving 88% accuracy.
- Pioneered and fine-tuned a multi-signal LLM classification model for SDG labeling using Meta LLaMA-3 (8B), evolving from few-shot to Chain-of-Thought reasoning, improving classification consistency by 30%+ and enabling interpretable, production-grade SDG alignment for large-scale nonprofit mission analysis.

**GATX Corporation**

**Jun 2025 - Dec 2025**

*Data Science Intern*

*Chicago, IL, USA*

- Led BI migration of 4 post-acquisition reports from Cognos to Power BI, re-architecting 20+ SQL queries and redesigning DirectQuery semantic models, reducing load times by 40% and increasing reporting accuracy by 30%. Partnered with post-acquisition business owners to scope and rebuild 4 mission-critical Cognos workflows in Power BI, re-architecting 20+ SQL queries and DirectQuery models, led reconciliation and UAT, and secured business sign-off with no material discrepancies while reducing load time by 40%.
- Owned end-to-end testing, data reconciliation, and validation for post-acquisition reports across Power BI, conducting UAT and securing business sign-off to ensure zero data discrepancy across merged systems.

**ZS Associates (Clients: AbbVie, Janssen, Gilead)**

**May 2022 - Aug 2024**

*Senior Data Engineer*

*Pune, India*

- Architected Spark-based ETL pipelines on AWS (Airflow, Athena, RDS) and Palantir Foundry with MDM integration, SCD-based historical tracking, pivoting logic and validation, improving pipeline efficiency by 35%.
- Designed and maintained a multi-tab, real-time Tableau reporting suite integrating Salesforce and CSV data sources to track patient experience KPIs (enrollment, service opt-ins, continuation/discontinuation), reducing report generation time by 40%.
- Spearheaded a scheduled SMS patient adherence pipeline, driving 30% opt-in growth across 100K+ active patients while maintaining HIPAA compliance across all concurrent Agile projects with 100% on-time delivery. Earned the Best Performance Award.
- Modeled patient-journey and business-entity datasets in Delta Lake and Snowflake using SCD Type 2 and temporal patterns, powering downstream analytics across 15+ stakeholder dashboards.

**S&P Global | LTI Mindtree**

**Jul 2020 - May 2022**

*Data Engineer*

*Mumbai, India*

- Developed scheduled ETL pipelines processing 2B+ records for cloud migration via Azure DevOps and Databricks, reducing system downtime by 30% and improving query performance by 45% through partitioning and indexing optimizations.
- Implemented data quality checks and reconciliation frameworks across batch pipelines to ensure 100% completeness and accuracy of migrated datasets, resolving data breaks through root-cause analysis.
- Collaborated with cross-functional teams to define data contracts, schema-evolution strategies, and CI/CD release workflows using Jenkins, enabling reliable production pipeline deployments and reducing release cycle time.

## PROJECTS

**NCAA Tournament Seed Prediction**

**Feb 2026 - Apr 2026**

- Solved and streamlined an end-to-end pipeline spanning data ingestion, feature engineering, predictive modeling, and Tableau dashboards on multi-season NCAA basketball data and bagged 1st Place in the NCAA Men's Analytics Challenge, 2026.

**Chronic Kidney Disease Prediction Pipeline**

**Feb 2025 - May 2025**

- Integrated Snowflake to manage and query 2M+ clinical records, automating data pipelines for preprocessing, feature engineering, and model input preparation.
- Restructured and optimized ML models (Logistic Regression, Random Forest, XGBoost) on clinical data, improving F1-score by 25% through SMOTE and feature engineering; deployed a Streamlit app with SHAP-based interpretability.

## TECHNICAL SKILLS

**Languages & Frameworks:** Python, SQL, Scala, PySpark, R

**Big Data and Cloud:** Apache Spark, Kafka, Hive, Hadoop MapReduce, Databricks, Delta Lake, AWS (S3, EMR, Glue, Lambda, Athena, Redshift, RDS), Azure (Databricks, DevOps, Data Factory), GCP (BigQuery), Snowflake

**Data Modeling & Quality:** Dimensional modeling, SCD Type 1/2, MDM, Schema Evolution, Data Reconciliation, dbt

**DevOps & Visualization:** EPIC, Fabric, Airflow, Jenkins, Palantir Foundry, CI/CD Pipelines, Git, JIRA, Docker, Tableau, Power BI, MongoDB, Neo4J

**Data Science and ML:** Feature Engineering, scikit-learn, PyTorch, Regression, Classification, LLMs, RAG, Agentic AI, HuggingFace

## CERTIFICATIONS

- AI Engineer for Data Scientists Associate: DataCamp (Feb 2026)